

Insurance cost prediction using PyTorch

August 11, 2021

```
[1]: import torch
import torchvision
import torch.nn as nn
import pandas as pd
import matplotlib.pyplot as plt
import torch.nn.functional as F
from torchvision.datasets.utils import download_url
from torch.utils.data import DataLoader, TensorDataset, random_split

#!conda install numpy pytorch torchvision cpuonly -c pytorch -y
```

```
[2]: #Step 1: Download and explore the data

DATA_FILENAME = "insurance.csv"
dataframe_raw = pd.read_csv(DATA_FILENAME)
dataframe_raw.head()

#We're going to do a slight customization of the data, so that you every
↳ participant receives a slightly different version of the dataset. Fill in
↳ your name below as a string (enter at least 5 characters)
```

```
[2]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
[3]: your_name = "Todor" # at least 5 characters

#The customize_dataset function will customize the dataset slightly using your
↳ name as a source of random numbers.
```

```
[4]: def customize_dataset(dataframe_raw, rand_str):
    dataframe = dataframe_raw.copy(deep=True)
    # drop some rows
```

```

dataframe = dataframe.sample(int(0.95*len(dataframe)),u
↪random_state=int(ord(rand_str[0])))
# scale input
dataframe.bmi = dataframe.bmi * ord(rand_str[1])/100.
# scale target
dataframe.charges = dataframe.charges * ord(rand_str[2])/100.
# drop column
if ord(rand_str[3]) % 2 == 1:
    dataframe = dataframe.drop(['region'], axis=1)
return dataframe

```

```

[5]: dataframe = customize_dataset(dataframe_raw, your_name)
dataframe.head()

```

```

[5]:      age    sex    bmi  children  smoker    charges
647    40  female  25.9407         3     no   8252.28430
1272   43    male  28.3272         5     no  14478.33015
880    22    male  38.6280         3     no   3443.06400
957    24    male  29.7369         1     no  12609.88702
1209   59    male  41.1810         1     no  12347.17200

```

```

[6]: #Let us answer some basic questions about the dataset.

```

```

[7]: #Q: How many rows does the dataset have?

```

```

num_rows = dataframe.shape[0]
print(num_rows)

```

1271

```

[8]: #Q: How many columns doe the dataset have

```

```

num_cols = dataframe.shape[1]
print(num_cols)

```

6

```

[9]: #Q: What are the column titles of the input variables? Hint: sex is one of them.
↪ List the columns that are not numbers.

```

```

input_cols = ['age', 'sex', 'bmi', 'children', 'smoker']

# a little fix of the code

input_cols = dataframe.columns[:5].tolist()
categorical_cols = dataframe.select_dtypes(exclude='number').columns.tolist()

```

```
[10]: #Q: What are the column titles of output/target variable(s)?
```

```
output_cols = ['charges']
```

```
[11]: print(dataframe['charges'].min())  
print(dataframe['charges'].max())
```

```
import seaborn as sns
```


```
plt.title("Various Charges")  
sns.distplot(dataframe.charges, kde=False);
```

1121.8739

63770.42801

C:\ProgramData\Miniconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).

```
warnings.warn(msg, FutureWarning)
```




```
[12]: #Q: (Optional) What is the minimum, maximum and average value of the charges_
      ↪column? Can you show the distribution of values in a graph? Use this data_
      ↪visualization cheatsheet
```

```
plt.title("Distribution of Charge")

sns.distplot(dataframe.charges)
plt.axvline(x=dataframe['charges'].min(),color="red")
plt.axvline(x=dataframe['charges'].max(),color="red")
plt.axvline(x=dataframe['charges'].mean(),color="red")
```

```
[12]: <matplotlib.lines.Line2D at 0x2868c5720a0>
```



```
[13]: #Step 2: Prepare the dataset for training
```

```
[14]: #We need to convert the data from the Pandas dataframe into a PyTorch tensors_
      ↪for training. To do this, the first step is to convert it numpy arrays. If_
      ↪you've filled out input_cols, categorical_cols and output_cols correctly,_
      ↪this following function will perform the conversion to numpy arrays.
```

```
[15]: def dataframe_to_arrays(dataframe):
      # Make a copy of the original dataframe
      dataframe1 = dataframe.copy(deep=True)
```

```

# Convert non-numeric categorical columns to numbers
for col in categorical_cols:
    dataframe1[col] = dataframe1[col].astype('category').cat.codes
# Extract input & outputs as numpy arrays
inputs_array = dataframe1[input_cols].to_numpy()
targets_array = dataframe1[output_cols].to_numpy()
return inputs_array, targets_array

```

[16]: *#Read through the Pandas documentation to understand how we're converting ↪ categorical variables into numbers.*

[17]: `inputs_array, targets_array = dataframe_to_arrays(dataframe)`
`inputs_array, targets_array`

[17]: (array([[40. , 0. , 25.9407, 3. , 0.],
 [43. , 1. , 28.3272, 5. , 0.],
 [22. , 1. , 38.628 , 3. , 0.],
 ...,
 [24. , 1. , 25.974 , 0. , 0.],
 [39. , 0. , 46.398 , 0. , 0.],
 [58. , 1. , 42.18 , 0. , 0.]]),
 array([[8252.2843],
 [14478.33015],
 [3443.064],
 ...,
 [1969.614],
 [5662.225],
 [11365.952]]))

[18]: *#Q: Convert the numpy arrays inputs_array and targets_array into PyTorch ↪ tensors. Make sure that the data type is torch.float32.*

```

inputs = torch.from_numpy(inputs_array).type(torch.float32)
targets = torch.from_numpy(targets_array).type(torch.float32)
inputs, targets

```

[18]: (tensor([[40.0000, 0.0000, 25.9407, 3.0000, 0.0000],
 [43.0000, 1.0000, 28.3272, 5.0000, 0.0000],
 [22.0000, 1.0000, 38.6280, 3.0000, 0.0000],
 ...,
 [24.0000, 1.0000, 25.9740, 0.0000, 0.0000],
 [39.0000, 0.0000, 46.3980, 0.0000, 0.0000],
 [58.0000, 1.0000, 42.1800, 0.0000, 0.0000]]),
 tensor([[8252.2842],
 [14478.3301],
 [3443.0640],
 ...,
 ...])

```
[ 1969.6140],  
[ 5662.2251],  
[11365.9521]])
```

```
[19]: inputs.dtype, targets.dtype
```

```
[19]: (torch.float32, torch.float32)
```

```
[20]: #Next, we need to create PyTorch datasets & data loaders for training &  
→validation. We'll start by creating a TensorDataset.
```

```
[21]: dataset = TensorDataset(inputs, targets)
```

```
[22]: ***Q: Pick a number between 0.1 and 0.2 to determine the fraction of data that  
→will be used for creating the validation set. Then use random_split to  
→create training & validation datasets. **
```

```
[23]: val_percent = 0.18 # between 0.1 and 0.2  
val_size = int(num_rows * val_percent)  
train_size = num_rows - val_size  
  
train_ds, val_ds = random_split(dataset, [train_size, val_size])  
# Use the random_split function to split dataset into 2 parts of the desired  
→length  
  
#Finally, we can create data loaders for training & validation.
```

```
[24]: #Q: Pick a batch size for the data loader.
```

```
[25]: batch_size = 25
```

```
[26]: train_loader = DataLoader(train_ds, batch_size, shuffle=True)  
val_loader = DataLoader(val_ds, batch_size)
```

```
[27]: #Let's look at a batch of data to verify everything is working fine so far.
```

```
[28]: for xb, yb in train_loader:  
    print("inputs:", xb)  
    print("targets:", yb)  
    break
```

```
inputs: tensor([[46.0000,  1.0000, 44.8162,  2.0000,  0.0000],  
                [27.0000,  1.0000, 36.1693,  3.0000,  0.0000],  
                [21.0000,  1.0000, 32.1623,  0.0000,  0.0000],  
                [24.0000,  0.0000, 26.8897,  0.0000,  0.0000],  
                [31.0000,  1.0000, 29.8424,  1.0000,  0.0000],
```

```

[19.0000, 1.0000, 33.9549, 0.0000, 0.0000],
[27.0000, 1.0000, 33.8550, 0.0000, 0.0000],
[35.0000, 1.0000, 27.1062, 3.0000, 1.0000],
[21.0000, 1.0000, 30.3696, 0.0000, 0.0000],
[19.0000, 0.0000, 30.9690, 0.0000, 1.0000],
[22.0000, 0.0000, 25.7298, 0.0000, 0.0000],
[48.0000, 0.0000, 35.8530, 2.0000, 0.0000],
[52.0000, 0.0000, 34.1658, 1.0000, 0.0000],
[18.0000, 0.0000, 32.3731, 0.0000, 0.0000],
[44.0000, 0.0000, 40.4928, 0.0000, 0.0000],
[40.0000, 0.0000, 33.0891, 1.0000, 0.0000],
[45.0000, 1.0000, 26.9841, 5.0000, 0.0000],
[45.0000, 1.0000, 44.1836, 0.0000, 0.0000],
[31.0000, 0.0000, 36.2748, 1.0000, 0.0000],
[26.0000, 0.0000, 25.0971, 0.0000, 0.0000],
[62.0000, 0.0000, 43.5120, 0.0000, 0.0000],
[50.0000, 1.0000, 29.3151, 0.0000, 0.0000],
[18.0000, 0.0000, 33.6385, 0.0000, 0.0000],
[56.0000, 1.0000, 35.2869, 2.0000, 1.0000],
[51.0000, 0.0000, 40.3873, 3.0000, 0.0000]])
targets: tensor([[ 8733.2295],
[ 4846.9199],
[ 1906.3583],
[ 2842.7607],
[ 4441.2134],
[ 1639.5631],
[ 2494.0220],
[19361.9980],
[ 2104.1133],
[16884.9238],
[ 2731.9121],
[10043.2490],
[10797.3359],
[ 7323.7349],
[12797.2100],
[ 6500.2358],
[ 9788.8662],
[ 7448.4038],
[ 4738.2681],
[ 3176.8159],
[13470.8604],
[ 8827.2100],
[ 2203.7358],
[43813.8672],
[11436.7383]])

```

[29]: *#Step 3: Create a Linear Regression Model*

```
[30]: #Our model itself is a fairly straightforward linear regression (we'll build  
→more complex models in the next assignment).
```

```
[31]: input_size = len(input_cols)  
output_size = len(output_cols)  
input_size, output_size
```

```
[31]: (5, 1)
```

```
[32]: #Q: Complete the class definition below by filling out the constructor  
→(.__init__), forward, training_step and validation_step methods.  
  
#Hint: Think carefully about picking a good loss fuction (it's not cross  
→entropy). Maybe try 2-3 of them and see which one works best. See https://  
→pytorch.org/docs/stable/nn.functional.html#loss-functions
```

```
[33]: class InsuranceModel(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.linear = nn.Linear(input_size, output_size) # fill this (hint: use  
→input_size & output_size defined above)  
  
    def forward(self, xb):  
        out = self.linear(xb) # fill this  
        return out  
  
    def training_step(self, batch):  
        inputs, targets = batch  
        # Generate predictions  
        out = self(inputs)  
        # Calcuete loss  
        loss = F.l1_loss (out, targets) # fill this  
        # mse_loss works best when the data is normalized. since this is  
→skewed, use l1_loss  
        return loss  
  
    def validation_step(self, batch):  
        inputs, targets = batch  
        # Generate predictions  
        out = self(inputs)  
        # Calculate loss  
        loss = F.l1_loss (out, targets) # fill this  
  
        →  
        return {'val_loss': loss.detach()}  
  
    def validation_epoch_end(self, outputs):  
        batch_losses = [x['val_loss'] for x in outputs]
```



```

        epoch_loss = torch.stack(batch_losses).mean()    # Combine losses
        return {'val_loss': epoch_loss.item()}

    def epoch_end(self, epoch, result, num_epochs):
        # Print result every 20th epoch
        if (epoch+1) % 20 == 0 or epoch == num_epochs-1:
            print("Epoch [{}], val_loss: {:.4f}".format(epoch+1,
→result['val_loss']))

```

[34]: *#Let us create a model using the InsuranceModel class. You may need to come
→back later and re-run the next cell to reinitialize the model, in case the
→loss becomes nan or infinity.*

[35]: `model = InsuranceModel()`

[36]: *#Let's check out the weights and biases of the model using model.parameters.*

[37]: `list(model.parameters())`

[37]: [Parameter containing:
tensor([[0.4234, 0.1851, 0.1789, -0.1294, 0.3414]], requires_grad=True),
Parameter containing:
tensor([-0.0034], requires_grad=True)]

[38]: *#Step 4: Train the model to fit the data*

[39]: *#To train our model, we'll use the same fit function explained in the lecture.
→That's the benefit of defining a generic training loop - you can use it for
→any problem.*

```

[40]: def evaluate(model, val_loader):
        outputs = [model.validation_step(batch) for batch in val_loader]
        return model.validation_epoch_end(outputs)

    def fit(epochs, lr, model, train_loader, val_loader, opt_func=torch.optim.SGD):
        history = []
        optimizer = opt_func(model.parameters(), lr)
        for epoch in range(epochs):
            # Training Phase
            for batch in train_loader:
                loss = model.training_step(batch)
                loss.backward()
                optimizer.step()
                optimizer.zero_grad()
            # Validation phase
            result = evaluate(model, val_loader)
            model.epoch_end(epoch, result, epochs)

```

```
        history.append(result)
    return history
```

[41]: *#Q: Use the evaluate function to calculate the loss on the validation set*
→ before training.

[42]: `result = evaluate(model, val_loader) # Use the the evaluate function`
`print(result)`

```
{'val_loss': 13696.3515625}
```

[43]: *#We are now ready to train the model. You may need to run the training loop*
→ many times, for different number of epochs and with different learning
→ rates, to get a good result. Also, if your loss becomes too large (or nan),
→ you may have to re-initialize the model by running the cell model =
→ InsuranceModel(). Experiment with this for a while, and try to get to as low
→ a loss as possible.

[44]: *#Q: Train the model 4-5 times with different learning rates & for different*
→ number of epochs.

#Hint: Vary learning rates by orders of 10 (e.g. 1e-2, 1e-3, 1e-4, 1e-5, 1e-6)
→ to figure out what works.

[45]: `epochs = 100`
`lr = 1e-1`
`history1 = fit(epochs, lr, model, train_loader, val_loader)`

```
Epoch [20], val_loss: 7099.5327
Epoch [40], val_loss: 7091.5654
Epoch [60], val_loss: 7086.2061
Epoch [80], val_loss: 7074.0610
Epoch [100], val_loss: 7069.6069
```

[46]: `epochs = 40`
`lr = 1e-2`
`history2 = fit(epochs, lr, model, train_loader, val_loader)`

```
Epoch [20], val_loss: 7063.9424
Epoch [40], val_loss: 7063.0249
```

[47]: `epochs = 50`
`lr = 1e-1`
`history3 = fit(epochs, lr, model, train_loader, val_loader)`

```
Epoch [20], val_loss: 7054.9111
Epoch [40], val_loss: 7048.6455
Epoch [50], val_loss: 7047.9463
```

```
[48]: epochs = 50
      lr = 1e-2
      history4 = fit(epochs, lr, model, train_loader, val_loader)
```

```
Epoch [20], val_loss: 7042.9736
Epoch [40], val_loss: 7042.3213
Epoch [50], val_loss: 7041.7905
```

```
[49]: epochs = 50
      lr = 1e-3
      history5 = fit(epochs, lr, model, train_loader, val_loader)
```

```
Epoch [20], val_loss: 7041.7705
Epoch [40], val_loss: 7041.7827
Epoch [50], val_loss: 7041.7588
```

```
[50]: #Q: What is the final validation loss of your model?
```

```
[51]: val_loss = history5[-1]['val_loss']
      val_loss
```

```
[51]: 7041.7587890625
```

```
[52]: #Step 5: Make predictions using the trained model
```

```
[53]: #Q: Complete the following function definition to make predictions on a single_
      ↪ input
```

```
[54]: def predict_single(input, target, model):
      inputs = input.unsqueeze(0)
      predictions = model(inputs)           # fill this
      prediction = predictions[0].detach()
      print("Input:", input)
      print("Target:", target)
      print("Prediction:", prediction)
```

```
[55]: input, target = val_ds[0]
      predict_single(input, target, model)
```

```
Input: tensor([50.0000,  1.0000, 28.0830,  0.0000,  0.0000])
Target: tensor([8442.6670])
Prediction: tensor([10508.6836])
```

```
[56]: input, target = val_ds[10]
      predict_single(input, target, model)
```

```
Input: tensor([61.0000,  0.0000, 34.5876,  0.0000,  0.0000])
Target: tensor([13429.0352])
```

Prediction: tensor([12854.3486])

```
[57]: input, target = val_ds[23]  
      predict_single(input, target, model)
```

Input: tensor([40.0000, 0.0000, 32.8560, 0.0000, 0.0000])

Target: tensor([5910.9438])

Prediction: tensor([7661.6782])